

# A Comprehensive Benchmark of Deep Learning Libraries on Mobile Devices

Qiyang Zhang, Xiang Li, Xiangying Che, Xiao Ma, Ao Zhou,  
Mengwei Xu, Shangguang Wang, Yun Ma, Xuanzhe Liu




# DL Inference on Smartphones

- Increasingly popular DL



Object Detection



NN accelerators

- Emerging DL libraries




# DL libraries are not fully understood




...


- Existing Benchmarks mainly focus on:



server



NN design



AI chip

Benchmark	Scenario	Benchmark objective	Supported DL libs
MLPerf [48]	T/I@S/E	Hardware	/
DeepBench [16]	T/I@S/E	Hardware	/
DAWNBench [15]	T/I@S	Hardware, algorithm, and DL libs	/
AI Matrix [7]	I@S	Hardware and DL libs	4
AI-Benchmark [34]	I@E	Hardware	1
Fathom [19]	T/I@S	Algorithm	1
gaugeNN [22]	I@E	DL apps and models	1
AIIA [13]	I@E	Hardware	3
This work	I@E	DL libs	6

The comparison of existing benchmarks and ours.

- A comprehensive benchmark for **on-device DL inference**

The benchmark triumphs at the aspect of rich support for more various **DL libs**

# Measurement Settings

- Rich support
  - the most popular DL libs: **TFLite, NCNN, MNN, PyTorch Mobile, Mace, SNPE.**
  - **15** models in total, for various tasks and different model precision.

# Measurement Settings

Supported DL libs and models in this work.

Models	Tasks	TFLite	ncnn	mnn	MACE	PyTorchMobile	SNPE
mobilenetV1 [29]	image classification	$C_{32,8}-G_{32,8}-D_8$	$C_{32,8}-G_{32,8}$	$C_{32,8}-G_{32,8}$	$C_{32,8}-G_{32}$	$C_{32,8}$	$C_{32,8}-G_{32,8}-D_8$
mobilenetV2 [55]	image classification	$C_{32,8}-G_{32,8}-D_8$	$C_{32,8}$	-	-	-	$C_{32,8}-G_{32,8}-D_8$
inceptionV3 [60]	image classification	$C_{32,8}-G_{32,8}-D_8$	$C_{32,8}$	-	-	-	$C_{32,8}-G_{32,8}-D_8$
inceptionV4 [59]	image classification	$C_{32,8}-G_{32,8}-D_8$	$C_{32,8}$	-	-	-	$C_{32,8}-G_{32,8}-D_8$
vgg16 [57]	image classification	$C_{32,8}-G_{32,8}-D_8$	$C_{32,8}$	-	-	-	$C_{32,8}-G_{32,8}-D_8$
squeezezenet [32]	image classification	$C_{32,8}-G_{32,8}-D_8$	$C_{32,8}-U_{32,8}$	$C_{32}-U_{32}$	$C_{32}-U_{32}$	$C_{32,8}$	$C_{32,8}-G_{32,8}-D_8$
nasnet_mobile [84]	image classification	$C_{32}-G_{32}$	-	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}$	-
densenet [31]	image classification	$C_{32}-G_{32}$	-	$C_{32}-G_{32}$	-	$C_{32}$	$C_{32}-G_{32}$
mnasnet [61]	image classification	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}$	$C_{32}-G_{32}$
resnetv2_50 [58]	image classification	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}$	$C_{32}-G_{32}$
deeplabv3 [81]	semantic segmentation	$C_{32}-G_{32}$	-	$C_{32}-G_{32}$	$C_{32}-G_{32}$	-	-
ssd_mobilenetV1 [46]	object detection	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}$	-
yolo-fastest [80]	object detection	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	-	-	-
yolo3 [53]	object detection	$C_{32}-G_{32}$	$C_{32}-G_{32}$	$C_{32}-G_{32}$	-	-	-
albert_tiny [70]	text classification	$C_{32}-G_{32}$	-	$C_{32}-G_{32}$	-	-	-

C/G/D: CPU/GPU/DSP  
32/8: FP32/INT8

# Measurement Settings

- Rich support
  - The most popular DL libs: **TFLite, NCNN, MNN, PyTorchMobile, Mace, SNPE.**
  - **15** models in total, for various tasks and different model precision.
- Devices
  - **10** different device models with various SoCs and GPUs
- Detailed metrics
  - **The inference time and operator-level information.**


# Research goal

**(The first) measurement** to understand how DL libs affect inference.

- **Performance Fragmentation**
- Impacts of **Quantization and Hardware**
- **Operator-level Integration**
- **Cold-start Inference**
- **Longitudinal Inference**



# Analysis Workflow



# Performance Fragmentation

Pytorch-M	MNN	TFLite	nncnn	SNPE	Mace					
MODEL	GP5	HE8	HM	MI11	MI9	MZ16	OP9	R9	RN9	S21
mobilenetV1	green	red	green	dark blue	red			green	green	red
mobilenetV2	green	red	green	red	red		green	green	green	red
inceptionV3	yellow	red	red	red	red		green	red		red
inceptionV4	yellow	red	green	red	red		green	green		red
vgg16	purple	purple	purple	purple	purple		purple	purple	purple	purple
squeezezenet	green	green	green	green	green		green	green	red	
mnasnet	green	green	dark blue	green	green		green	green	green	
resnetV2_50	dark blue		dark blue	dark blue	dark blue	dark blue				
nasnet_mobile	dark blue		dark blue	dark blue	dark blue	dark blue				
densenet	yellow	red	red	red	red	green	red	red	red	
ssd_mobilenetV1	dark blue		dark blue	dark blue	dark blue	dark blue				
deeplabV3	dark blue		dark blue	dark blue	dark blue	dark blue				
yolo-fastest	green	red	red	red	red	green	green	green	red	
yolo3	red	red	red	red	red	green	red	red	red	
albert_tiny	dark blue	red	red	red	red	dark blue	dark blue	dark blue	red	red
mobilenetV1_INT8	dark blue	dark blue	green	dark blue	dark blue		dark blue	dark blue	dark blue	dark blue
mobilenetV2_INT8	dark blue		dark blue	dark blue	dark blue	dark blue				
inceptionV3_INT8	purple	dark blue	dark blue	dark blue	dark blue		dark blue	dark blue	dark blue	dark blue
inceptionV4_INT8	dark blue	green	dark blue	dark blue	dark blue	dark blue				
squeezezenet_INT8	green	green	green	green	green		green	green	green	green
vgg16_INT8	purple	purple	purple	dark blue	purple		purple	purple	purple	purple

Best-performing lib on CPU

**DL lib with the  
smallest inference  
time in 6 DL libs**



**The DL lib with fastest speed  
to run "albert\_tiny" on  
"Samsung S21" is "MNN".**

# Performance Fragmentation

	MNN	V/G/C	TFLite	ncnn	SNPE	Mace				
MODEL	GP5	HE8	HM	MI11	MI9	MZ16	OP9	R9	RN9	S21
mobilenetV1			C	C			C			
mobilenetV2			C	C		C	C			C
inceptionV3			C	C			C		V	V
inceptionV4			C	C			C		V	V
vgg16			C	C		C	C			
squeezezenet				C	C		V			C
mnasnet										
resnetV2_50										
nasnet_mobile										
densenet										
ssd_mobilenetV1										
deeplabV3										
yolo-fastest		G	V	G	C	G	C		V	G
yolo3	V	G	G	C	G	C	C	G	V	G
albert_tiny	0	0	V	G	G	G				G
mobilenetV1_INT8										
mobilenetV2_INT8							G			
inceptionV3_INT8										
inceptionV4_INT8										
squeezezenet_INT8										
vgg16_INT8			V	V	C			V	V	

Best-performing lib on CPU/GPU

**There is no one-size-fit-all DL lib  
for optimal performance across  
models and hardware.**

# Performance Fragmentation


Performance gap of DL libs can be large.

Models	Best vs. Worst		Best vs. 2nd Best		
	CPU (x)	GPU (x)	CPU (x)	GPU (x)	
mobilenetV1	4.0~15.4 (8.7)	1.7~14.	the longer one divided by the shorter one	.5)	1.0~4.0 (1.9)
mobilenetV2	5.6~18.8 (11.2)	2.9~15.		.5)	1.0~2.9 (1.6)
inceptionV3	2.6~5.6 (3.8)	3.0~13.4 (1.1)	1.1~2.4 (1.7)	1.0~4.0 (2.1)	
inceptionV4	2.0~5.4 (3.2)	2.4~11.0 (5.8)	1.1~2.0 (1.5)	1.0~3.6 (2.0)	
vgg16	7.1~54.3 (16.2)	4.4~7.0 (5.5)	1.3~4.2 (2.4)	1.1~2.2 (1.5)	
squeezezenet	4.6~19.9 (9.1)	1.9~12.6 (5.9)	1.0~5.9 (2.5)	1.1~2.5 (1.6)	
average	<u>8.7</u>	<u>6.0</u>	<u>1.9</u>	<u>1.8</u>	

The performance gaps of different DL libs

# Performance Fragmentation


**With software heterogeneity, the model structure is not the sole factor that determines relative performance.**



***Implication:*** To pursue the optimal performance, the developers need to incorporate different DL libs and dynamically load one based on the current model and hardware platform.

# Performance Fragmentation

Benefit brought by INT8 quantization is under expectation.




0.8×–3.0× faster  
than FLOAT

Best inference speed across all DL libs

# Performance Fragmentation

Benefit brought by INT8 quantization is under expectation.




0.8×–3.0× faster  
than FLOAT

**Implication:** There exists great potential at software level to accelerate the inference of quantized models.

# Impacts of Hardware

**GPU can not always accelerate DL inference.**




speedup 1.4×–  
1.9× compared to  
CPU

Best inference speed across all DL libs

# Impacts of Hardware


On INT8-based models, GPU can hardly bring any benefit.



**Implication:** Observations motivate developers to focus on GPU optimization. It also motivates researchers to design models suitable for GPU.

# Impacts of Hardware

DSP can significantly accelerate INT8 model in most cases.



reduce inference  
time of INT8 model  
by **2.0×–12.9×**

**Implication:** The current DL libs can not fully exploit the capacity of each hardware.

# Operator-level Integration



**How about integrating the best-performing operator from DL libs?**

Oracle lib that combines the fastest operator from those DL libs

Models	Mace	tflite	SNPE	ncnn	Oracle time
mobilenetV1				14.4	13.5 (↓6.1%)
mobilenetV2				14.4	10.6 (↓26.3%)
inceptionV3				123	86.3 (↓29.9%)
inceptionV4				74.9	180.3 (↓7.6%)
vgg16	180.3	73.1	341.7	409.0	73.1 (↓0%)


The benefits that integrate the wisdom of DL libs ( ms )

**Oracle time brings  
inference time reduction**


**Implication:** Those diversities need to be unified before the operator implementation can be combined.

# Cold-start Inference

The first inference beginning from model loading




# Breakdown of Cold-start Inference



**Memory preparation** contributes to the largest overhead in cold-start inference.



# Breakdown of Cold-start Inference



***Implication:*** Potential solutions include speeding up memory preparation using multiple threads and generating pipeline to run model loading memory preparation and inference simultaneously.

# Longitudinal Analysis

↓  
Better




Time →

**The inference performance evolvement across time of TFLite**

The performance of DL libs are continuously improving in early years, but becomes relatively stable since 2020.


# Longitudinal Analysis



# Longitudinal Analysis



Better



**Implication:** The current open-source ecosystems is possibly due to a comprehensive benchmarking tool available for developers to test commits.

# Summary

- A comprehensive benchmark to quantitatively understand inference performance of DL libs.
- Lead to insightful implications for complete landscape of DL libs ecosystem.

Please check benchmark at

➤ <https://github.com/UbiquitousLearning/MobileDLFrameworksBenchmark>

Thanks for your attention!

